

ACM SIG AI - SEM 2 Recruitment Tasks

Task 0: Complete the following courses

To ensure you have the foundational knowledge needed for the upcoming tasks, please complete the following courses:

1. **Kaggle: Intro to Machine Learning** - [Kaggle Course](#)
2. **Kaggle: Supervised Machine Learning** - [Kaggle Supervised Learning](#)
3. **Kaggle: Unsupervised Learning** - [Kaggle Unsupervised Learning](#)
4. **Scikit-Learn Official Documentation** - [Scikit-Learn Docs](#)

Note: It's highly recommended that you complete these courses before proceeding with the tasks.

Task 1: Implement a Regression Model (Building from Scratch)

Objective: Select a dataset that requires substantial preprocessing, and implement a Linear Regression model from scratch. You are required to handle preprocessing techniques and build an end-to-end workflow. This task evaluates your understanding of regression, data preprocessing, and model evaluation.

1. **Dataset:** Choose a dataset that requires extensive preprocessing.
 - Example: House prices, medical data, etc.
 - You should apply techniques like handling missing values, feature scaling, encoding categorical variables, etc.
2. **Preprocessing:** Perform necessary preprocessing steps:
 - Handle missing data.
 - Normalize/standardize features if necessary.
 - Visualize data and check for correlations using scatter plots, etc.
3. **Model:** Implement Linear Regression from scratch (do not use any inbuilt libraries like scikit-learn for the model).
 - Use basic linear algebra to perform calculations.

- Implement a simple gradient descent approach to optimize the model.
4. **Evaluation:**

- Evaluate your model using **Mean Squared Error (MSE)** and **Root Mean Squared Error (RMSE)**.
 - Perform a comparison between **Linear Regression** and **Polynomial Regression**. Plot the performance differences and analyze them.
-

Task 2: Dealing with Images (Classification)

Objective: Work with image data to build a classification model. You will explore image preprocessing, feature extraction, and model evaluation using different classifiers.

1. **Dataset:** Choose a small-scale image classification dataset (e.g., MNIST, CIFAR-10, Fashion-MNIST).
 - Download one of the mentioned datasets from Kaggle or other repositories.
 2. **Preprocessing:** Perform image preprocessing:
 - Normalize the pixel values (e.g., to the range [0, 1]).
 - Resize images to a fixed shape (if required).
 - Optionally, perform augmentation (e.g., rotations, flips).
 3. **Model:** Select a classification algorithm:
 - Start with a basic classifier such as **Logistic Regression**, **SVM**, or **Decision Trees**.
 - Optionally, implement a **simple Neural Network** if you are familiar with it.
 4. **Evaluation:**
 - Measure model performance using **accuracy**, **precision**, **recall**, and **F1-score**.
 - Generate and analyze a **confusion matrix**.
 - Visualize misclassified examples and correct ones.
-

Task 3: Text Analysis (Unsupervised Learning)

Objective: Work with text data to perform unsupervised learning through clustering. You will preprocess the text, extract features, and apply a clustering algorithm.

1. **Dataset:** Choose a small text dataset (e.g., product reviews, news headlines, or a collection of short articles).
 - A sample dataset could be found on Kaggle or any public text repository.
 2. **Preprocessing:**
 - Convert all text to lowercase.
 - Remove punctuation and special characters.
 - Optionally, remove stop words and apply stemming/lemmatization.
 3. **Model:**
 - Use **TF-IDF (Term Frequency-Inverse Document Frequency)** to convert the text into numerical format.
 - Apply **K-Means clustering** to group the text into 2 clusters.
 4. **Interpretation:**
 - Identify key terms for each cluster.
 - Explain why these terms belong to the same group.
 - **Bonus:** Calculate cosine similarity to find the two most similar texts in the dataset.
-

Task 4: Dealing with PyTorch

Objective: Understand basic PyTorch operations like tensor manipulation, matrix operations, and automatic differentiation using autograd. Complete the following steps to gain hands-on experience.

1. **Matrix Multiplication and Autograd:**
 - Create two random tensors of shape **3x3**.
 - Perform matrix multiplication between the two tensors.
 - Use PyTorch's **autograd** to compute the gradient of the result with respect to one of the input tensors.

- **Explain:** What happens in each step and how autograd computes gradients.
 - 2. **Broadcasting:**
 - Create a **3x1** tensor and a **1x3** tensor.
 - Use broadcasting to add the two tensors together.
 - Multiply the result by another **3x3** tensor.
 - **Explain:** What broadcasting is and how it works in this case.
 - 3. **Reshaping and Slicing:**
 - Create a tensor of shape **(6, 4)** using random values and reshape it to shape **(3, 8)**.
 - After reshaping, extract a slice (e.g., the first two columns, all rows).
 - **Explain:** What reshaping does and how slicing helps to extract specific parts of the tensor.
-

Submission Guidelines:

1. **Code:** All code should be written in Python. Please use Jupyter notebooks or Python scripts to submit your work.
2. **Documentation:** Provide clear explanations for each step of your process, including how you approached preprocessing, model selection, and evaluation.
3. **Visualization:** Ensure that you include relevant data visualizations (graphs, plots, etc.) where necessary.
4. **GitHub:** All completed tasks must be uploaded to a GitHub repository. The repository should include:
 - Make it **Public** with repo name as “<yourname>_ACM_SIGAI_Recr”
 - A clear README file explaining how to run the code and the results.
 - Separate folders for each task (Task 1, Task 2, etc.), with respective files clearly named.

Please ensure that the notebook or code is well-commented and easy to understand.
